

Smartphone Market Analysis & Forecasting

How Xiaomi lost its crown in India's smartphone market to Samsung & Apple (2020–2030)

Polished notebook prepared for GitHub / Resume — includes EDA, insights, and forecasting (Prophet).

1. Introduction

Context: The Indian smartphone market between 2020 and 2024 experienced a shift: Xiaomi — once market leader — lost ground to Samsung and Apple.

Objective:

- Analyze historical sales & revenue trends (2020–2024).
- Explain why Xiaomi lost its crown.
- Forecast brand sales and market share through 2030 (quantitative outlook).

Notes: This notebook expects the dataset file `smartphone_sales_updated_new.xlsx` to be in the same directory or the `/mnt/data/` path.

```
# 2. Setup & Imports
import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import plotly.express as px

df=pd.read_csv('smartphone_sales_updated.csv')
df
```

	Phone_ID	Brand	RAM_GB	Storage_GB	Screen_Size	Price_USD	\
0	1	OnePlus	8	256	7.15	395.00	
1	2	Google	12	256	6.44	228.00	
2	3	Xiaomi	16	1024	7.23	814.50	
3	4	Google	16	512	6.49	NaN	
4	5	Google	16	128	7.03	NaN	
...
1995	1996	Google	12	64	5.58	1,035.00	
1996	1997	Samsung	6	512	7.11	719.40	
1997	1998	Google	8	64	6.82	1,338.00	
1998	1999	Xiaomi	16	1024	7.45	460.80	
1999	2000	Apple	4	64	5.39	1,460.40	

	Battery_mAh	OS	Rating	Quantity_Sold	Sales_Revenue
Profit \ 6,517.50	3,227.00	Android	4.30	66	26,070.00
5,700.00	5,881.00	Android	3.50	100	22,800.00
112,220.00	3,857.00	iOS	4.80	496	NaN
97,295.00	4,268.00	iOS	4.70	319	389,180.00
4	5,865.00	Android	4.00	364	78,260.00
NaN
1995 111,003.75	5,893.00	Android	3.20	429	444,015.00
1996 22,726.50	3,943.00	Android	4.30	139	90,906.00
1997 93,994.50	5,326.00	Android	3.00	281	375,978.00
1998 47,616.00	5,429.00	Android	3.00	372	190,464.00
1999 93,404.75	5,060.00	iOS	3.80	307	373,619.00
	Year				
0	2023				
1	2023				
2	2023				
3	2025				
4	2023				
...	...				
1995	2023				
1996	2025				
1997	2022				
1998	2020				
1999	2020				

[2000 rows x 13 columns]

4. Data Overview & Cleaning

- Check for nulls, duplicates, dtypes.
- Basic cleaning performed: median imputation for numeric columns if required, ensure Sales_Revenue exists.

```
# Basic data checks
display(df.info())
```

```
print('\nNull counts:', df.isnull().sum())
print('\nDuplicate rows:', df.duplicated().sum())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Phone_ID        2000 non-null    int64  
 1   Brand            2000 non-null    object  
 2   RAM_GB           2000 non-null    int64  
 3   Storage_GB       2000 non-null    int64  
 4   Screen_Size      2000 non-null    float64 
 5   Price_USD        1900 non-null    float64 
 6   Battery_mAh     1940 non-null    float64 
 7   OS               2000 non-null    object  
 8   Rating           2000 non-null    float64 
 9   Quantity_Sold    2000 non-null    int64  
 10  Sales_Revenue    1900 non-null    float64 
 11  Profit           1940 non-null    float64 
 12  Year              2000 non-null    int64  
dtypes: float64(6), int64(5), object(2)
memory usage: 203.3+ KB
```

None

```
Null counts:
Phone_ID          0
Brand             0
RAM_GB            0
Storage_GB         0
Screen_Size        0
Price_USD         100
Battery_mAh       60
OS                0
Rating            0
Quantity_Sold     0
Sales_Revenue     100
Profit            60
Year              0
dtype: int64
```

Duplicate rows: 0

```
pd.crosstab(df['Brand'], df['OS'])
```

OS	Android	iOS
Brand		
Apple	277	52
Google	276	61

```

Huawei      277  53
OnePlus     265  61
Samsung     288  64
Xiaomi      272  54

# Identify invalid OS records
invalid_os = df[
    (df['Brand'].str.lower() != 'apple') & (df['OS'] == 'iOS')
]

print(f"Invalid OS records found: {invalid_os.shape[0]}")

Invalid OS records found: 293

# Correct invalid OS values
df.loc[
    (df['Brand'].str.lower() != 'apple') & (df['OS'] == 'iOS'),
    'OS'
] = 'Android'

df.loc[
    df['Brand'].str.lower() == 'apple',
    'OS'
] = 'iOS'

pd.crosstab(df['Brand'], df['OS'])

OS      Android  iOS
Brand
Apple      0   329
Google     337   0
Huawei     330   0
OnePlus    326   0
Samsung    352   0
Xiaomi     326   0

df

```

	Phone_ID	Brand	RAM_GB	Storage_GB	Screen_Size	Price_USD	\
0	1	OnePlus	8	256	7.15	395.00	
1	2	Google	12	256	6.44	228.00	
2	3	Xiaomi	16	1024	7.23	814.50	
3	4	Google	16	512	6.49	NaN	
4	5	Google	16	128	7.03	NaN	
...
1995	1996	Google	12	64	5.58	1,035.00	
1996	1997	Samsung	6	512	7.11	719.40	
1997	1998	Google	8	64	6.82	1,338.00	
1998	1999	Xiaomi	16	1024	7.45	460.80	
1999	2000	Apple	4	64	5.39	1,460.40	

	Battery_mAh	OS	Rating	Quantity_Sold	Sales_Revenue
Profit \ 0	3,227.00	Android	4.30	66	26,070.00
6,517.50					
1	5,881.00	Android	3.50	100	22,800.00
5,700.00					
2	3,857.00	Android	4.80	496	NaN
112,220.00					
3	4,268.00	Android	4.70	319	389,180.00
97,295.00					
4	5,865.00	Android	4.00	364	78,260.00
NaN					
...
...					
1995	5,893.00	Android	3.20	429	444,015.00
111,003.75					
1996	3,943.00	Android	4.30	139	90,906.00
22,726.50					
1997	5,326.00	Android	3.00	281	375,978.00
93,994.50					
1998	5,429.00	Android	3.00	372	190,464.00
47,616.00					
1999	5,060.00	iOS	3.80	307	373,619.00
93,404.75					
	Year				
0	2023				
1	2023				
2	2023				
3	2025				
4	2023				
...	...				
1995	2023				
1996	2025				
1997	2022				
1998	2020				
1999	2020				
[2000 rows x 13 columns]					
df.isnull().sum()					
Phone_ID	0				
Brand	0				
RAM_GB	0				
Storage_GB	0				
Screen_Size	0				
Price_USD	100				
Battery_mAh	60				
OS	0				

```

Rating          0
Quantity_Sold   0
Sales_Revenue   100
Profit          60
Year            0
dtype: int64

# Fill nulls for all numeric columns
for col in num_cols:
    median_value = df[col].median()
    df[col] = df[col].fillna(median_value)

df.isnull().sum()

Phone_ID        0
Brand           0
RAM_GB          0
Storage_GB      0
Screen_Size     0
Price_USD       0
Battery_mAh     0
OS              0
Rating          0
Quantity_Sold   0
Sales_Revenue   0
Profit          0
Year            0
dtype: int64

```

5. Feature Engineering

- Create Price_Segment from Price_USD.
- Aggregate brand-year level data for EDA and forecasting.

```

# Price segment function & application
def price_segment(price):
    if price < 100: return 'Budget'
    elif price < 400: return 'Mid'
    elif price < 600: return 'Mid-Premium'
    elif price < 800: return 'Premium'
    else: return 'Super-Premium'

if 'Price_Segment' not in df.columns:
    df['Price_Segment'] = df['Price_USD'].apply(price_segment)

display(df[['Brand', 'Price_USD', 'Price_Segment']].head())

```

	Brand	Price_USD	Price_Segment
0	OnePlus	395.00	Mid
1	Google	228.00	Mid
2	Xiaomi	814.50	Super-Premium

```

3   Google     887.30  Super-Premium
4   Google     887.30  Super-Premium

df.to_excel("smartphone_sales_clean.xlsx", index=False)

# Aggregate to brand-year level
year_brand = df.groupby(['Year', 'Brand'], as_index=False)
['Sales_Revenue'].sum()
brand_totals = df.groupby('Brand', as_index=False)
['Sales_Revenue'].sum().sort_values('Sales_Revenue', ascending=False)
display(brand_totals.head())

      Brand Sales_Revenue
4  Samsung  82,894,046.00
2  Huawei   78,718,611.00
1  Google    78,133,669.50
3  OnePlus   77,508,820.50
5  Xiaomi    76,310,108.50

```

6. Exploratory Data Analysis (EDA)

Key visualizations: revenue trend per brand, market share pie chart for top brands, price distribution, correlation heatmap, and scatter Price vs Quantity.

```

top_3_brands = (
    df.groupby('Brand')['Sales_Revenue']
    .sum()
    .sort_values(ascending=False)
    .head(3)
)

top_3_brands

Brand
Samsung    82,894,046.00
Huawei     78,718,611.00
Google      78,133,669.50
Name: Sales_Revenue, dtype: float64

top_brands_list = top_3_brands.index.tolist()
top_brands_df = df[df['Brand'].isin(top_brands_list)]

brand_revenue_trend = (
    top_brands_df
    .groupby(['Year', 'Brand'])['Sales_Revenue']
    .sum()
    .reset_index()
)
brand_revenue_trend.head()

```

```

      Year   Brand Sales_Revenue
0  2020  Google 11,897,164.50
1  2020  Huawei 11,738,408.50
2  2020  Samsung 14,171,061.50
3  2021  Google 14,418,819.50
4  2021  Huawei 16,207,841.00

```


```

import plotly.express as px

fig = px.line(
    brand_revenue_trend,
    x='Year',
    y='Sales_Revenue',
    color='Brand',
    markers=True,
    title='Revenue Trend Over Time – Top 3 Smartphone Brands'
)

fig.show()

```



"The top 3 brands show consistent revenue growth over time. One brand dominates due to higher pricing strategy, while others rely on volume sales. This indicates different business models within the same market."

```

fig = px.pie(
    brand_revenue_trend,
    names='Brand',
    values='Sales_Revenue',
    title='Market Share by Revenue (Top 3 Brands: 2020–2024)',
    hole=0.4  # optional → donut chart (looks professional)
)

fig.show()

```

Market Share by Revenue (Top 3 Brands: 2020–2024)




```
import plotly.express as px

fig = px.box(
    top_brands_df,
    x='Brand',
    y='Price_USD',
    title='Price Distribution of Top 3 Smartphone Brands',
    points='outliers' # shows extreme pricing
)


fig.show()
```

Price Distribution of Top 3 Smartphone Brands



```
# Correlation heatmap (numeric columns)
num_cols =
['RAM_GB', 'Storage_GB', 'Screen_Size', 'Price_USD', 'Battery_mAh', 'Rating',
 'Quantity_Sold', 'Sales_Revenue', 'Profit']
plt.figure(figsize=(10,7))
sns.heatmap(
    df[num_cols].corr(),
    annot=True,
    fmt='.2f',
    cmap='coolwarm',
    linewidths=0.5
)
```

```
plt.title('Correlation Heatmap of Smartphone Features & Sales Metrics')
plt.show()
```



"The heatmap shows that revenue and profit are strongly correlated with price and quantity sold, while technical features like RAM, battery, and screen size have negligible correlation. This suggests that consumer purchasing decisions are influenced more by pricing and brand positioning than raw specifications"

```
# Scatter: Price vs Quantity Sold (bubble-size by Sales_Revenue)
fig = px.scatter(
    top_brands_df,
    x='Price_USD',
    y='Quantity_Sold',
    size='Sales_Revenue',
    color='Brand',
    hover_data=['RAM_GB', 'Storage_GB'],
```

```

        title='Price vs Quantity Sold with Revenue Impact (Top 3 Brands)' ,
        size_max=50
    )
fig.show()

```

Price vs Quantity Sold with Revenue Impact (Top 3 Brands)



7. Brand-level Summary Metrics

Compute average price, average rating, models count, and total revenue for the three favorite brands.


```

model_count = (
    top_brands_df
    .groupby('Brand')[ 'Phone_ID' ]
    .nunique()
    .reset_index(name='Models_Count')
)

fig = px.bar(
    model_count,
    x='Brand',
    y='Models_Count',
    title='Product Portfolio Size – Top 3 Brands',
    text_auto=True
)
fig.show()

```


Product Portfolio Size – Top 3 Brands



```
avg_rating = (
    top_brands_df
    .groupby('Brand')['Rating']
    .mean()
    .reset_index()
)

fig = px.bar(
    avg_rating,
    x='Brand',
    y='Rating',
    title='Average Customer Rating – Top 3 Brands',
    text_auto='.2f'
)
fig.show()
```




Average Customer Rating – Top 3 Brands




"The EDA reveals that smartphone market success is driven primarily by pricing strategy and demand volume rather than hardware specifications. Top brands differentiate themselves through premium positioning, product diversity, or value pricing, while maintaining similar customer satisfaction levels."

```
import matplotlib.pyplot as plt
import seaborn as sns
```


```
for col in num_cols:  
    plt.figure(figsize=(6,2))  
    sns.boxplot(x=df[col])  
    plt.title(col)  
    plt.show()
```




Price_USD






Battery_mAh



Rating





```
num_cols = df.select_dtypes(include=['int64', 'float64']).columns

for col in num_cols:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5*IQR
```

```

upper_bound = Q3 + 1.5*IQR
median = df[col].median()

# Replace outliers using loc
df.loc[(df[col] < lower_bound) | (df[col] > upper_bound), col] =
median



print("Outliers replaced with median.")

Outliers replaced with median.


import matplotlib.pyplot as plt
import seaborn as sns

for col in num_cols:
    plt.figure(figsize=(6,2))
    sns.boxplot(x=df[col])
    plt.title(col)
    plt.show()


```




Storage_GB




Screen_Size




Price_USD




Battery_mAh






Rating



Quantity_Sold





why i choose linear reg model

"After performing EDA, I found that Price and Quantity Sold had the strongest correlation with Sales Revenue, while other features like RAM, battery, and screen size had minimal impact. Since the relationship was mostly linear and the dataset was structured, Linear Regression was the most suitable model.

I avoided complex machine learning models because they reduce interpretability. Linear Regression allows me to directly quantify how much revenue changes with price or sales volume, which is far more valuable for business decision-making."

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_absolute_error

X = df[['Price_USD', 'Quantity_Sold']]
y = df['Sales_Revenue']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

model = LinearRegression()
model.fit(X_train, y_train)

LinearRegression()

y_pred = model.predict(X_test)

r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)


r2, mae
(0.7867500299246708, 45961.111499194725)

coef_df = pd.DataFrame({
    'Feature': X.columns,
    'Coefficient': model.coef_
})

coef_df
      Feature  Coefficient
0  Price_USD        209.76
1  Quantity_Sold       730.77

import matplotlib.pyplot as plt

plt.scatter(y_test, y_pred)
plt.xlabel("Actual Revenue")
plt.ylabel("Predicted Revenue")
plt.title("Actual vs Predicted Revenue")
plt.show()
```



Modeling Summary

A multiple linear regression model was built to validate key EDA insights. The model confirms that pricing and quantity sold are the strongest predictors of sales revenue, while other product specifications contribute minimally. The simplicity and interpretability of the model make it suitable for business decision-making.